

SHREYASH PATANGE

Pune, Maharashtra

☎ +91-9130276761 ✉ shreyashpatange@gmail.com 🌐 GitHub in LinkedIn 📄 Leetcode 📁 Portfolio

PROFESSIONAL SUMMARY

Passionate about backend engineering, AI agents, and building systems that actually work in production. Quick to pick up new tools and frameworks, I enjoy digging into hard infrastructure problems and turning them into simple, reliable solutions. Driven by real impact, I like building intelligent systems that combine automation, AI, and clean engineering to save people time.

EDUCATION

Pune Institute Of Computer Technology
Electronics and Telecommunication

08-2022 – 07-2026
Pune, India

EXPERIENCE

Trainee Software Engineer
Intangles.AI

AUG 2026 – Present
Pune, India

- Converted from intern to full-time; building an **AI marketing agent** that automates the **end-to-end SDR workflow** – lead sourcing, qualification, and personalized outreach – replacing a manual prospecting process with an **LLM-driven** pipeline.
- Integrated **Apollo** for lead enrichment and **CRM** sync with automated **email and LinkedIn outreach sequencing**, moving qualified leads from discovery to first touch without manual handoff – cutting per-lead effort by **90%**.
- Designed **LLM-based lead scoring** and message personalization driven by prospect and company signals, enabling the marketing team to run outreach at scale and lifting client conversions by **50%**.

Backend Developer Intern
Intangles.AI

FEB 2026 – AUG 2026
Pune, India

- Built and deployed **HUSH**, an **autonomous SRE agent** on **AWS (ap-south-1)**, eliminating manual on-call investigation – upon a **PagerDuty** alert, the system autonomously queries **RabbitMQ**, **Kubernetes** application logs, **Dash0**, and **Prometheus**, identifies the root cause, and posts a structured diagnosis and fix to **Slack** in under **2 minutes**, reducing investigation time from **30 minutes to 5 minutes** and targeting **50%+ fewer on-call pages**.
- Designed a **3-tier runbook engine** with Kubernetes safety guards across **5 environments**; engineered **CrewAI** multi-agent fuel theft detection achieving **97% accuracy** and a deterministic **OpenClaw** pipeline with **60/65** validation points.
- Integrated **AWS Bedrock** and **LLM-based reasoning** to enable an extensible agentic architecture where new investigative agents can be added per alert type, allowing the system to continuously expand its coverage without human re-engineering.

Product Engineer Intern
Performena

AUG 2025 – DEC 2025
Remote

- Built an AI-powered Market & Media Research Analyzer generating detailed company reports by combining financial, media, and digital data – cutting report generation from **2 hours to under 4 minutes (98% reduction)**.
- Helped businesses identify market trends, competitor strategies, and investment opportunities through fully automated analysis, enabling teams to **work efficiently at scale**.
- Delivered data-driven reports supporting strategic planning and faster decision-making, eliminating hours of manual research per client engagement.

PROJECTS

CogniLearn Skilling Platform 🌐 | React, Flask, OpenAI API, GitHub API, Transformers

MAR 2025

- Implemented an AI-driven learning portal offering personalized skill roadmaps, curated notes, and recommended courses for engineering students.
- Integrated GitHub API and Gemini AI to recommend skill-aligned project summaries and guide open-source contributions.
- Developed an AI interview system that generates questions from user skills and evaluates voice responses using Whisper and LLaMA 3.3-70B.

Distributed Rate Limiter & API Gateway 🌐 | Node.js, TypeScript, Redis, Prometheus, k6

JUL 2026

- Built an API gateway enforcing correct per-client rate limits across **multiple instances** behind a load balancer via atomic **Redis Lua scripts** (EVALSHA) – verified with a test proving **2 concurrent gateway processes** sharing one Redis allow exactly the configured limit, not double it.
- Benchmarked **three algorithms** (token bucket, sliding window, fixed window); empirically proved fixed window’s boundary-burst flaw – **10 requests allowed against a limit of 5**, vs. sliding window correctly holding at **5**.
- Load-tested with **k6**, instrumented with **Prometheus/Grafana**, and deployed to production (Render) with a fail-open design so Redis outages don’t take down the upstream.

TECHNICAL SKILLS

Languages: Python, C, C++, JavaScript, SQL
Frontend Technology: HTML5, CSS3, React.js, Tailwind CSS, Streamlit
Backend Technology: Node.js, Express.js, RESTful APIs
ML Libraries: Scikit-learn, Keras, Pandas, NumPy, Matplotlib, Seaborn
AI/Agentic Tools: CrewAI, OpenClaw, LLM Integration, AWS Bedrock, Apollo, CRM Automation
Infrastructure: Kubernetes, Docker, Redis, RabbitMQ, PagerDuty, Dash0, Prometheus
Databases: MongoDB, MySQL, PostgreSQL

ACHIEVEMENTS

- **Springer LNNS (Scopus-indexed)** – “ML Approaches for Symptom-Based Syndrome Analysis Using SVC”
- Paper submitted to **PEIS 2026** – “Design and Implementation of Advanced Tool for Media Strategy and Planning”
- Runner-up at **INSPIRON 4.0 National Hackathon**, COEP

CERTIFICATIONS

- Supervised Machine Learning: Regression and Classification
- MITU Skillgies: Data Science with Machine Learning